



Sail Optimization Engine V2: A Mathematical Framework for Spike-Aware Yield Optimization Across Networks and Stablecoins

Alvaro Alonso Miguel
Senior Executive Engineer, Sail

February 2026

Abstract

We present the second-generation optimization engine for **Sail**'s autonomous cross-network portfolio management system, extending the theoretical framework established in our previous work [13] with a novel *Smart Spike Detection Framework*. The engine addresses the fundamental challenge of distinguishing transient yield spikes from sustainable trends via a z-score based statistical classifier, combined with an aggressive profitability gatekeeper enforcing tier-based execution cycles (Silver: 6h, Gold: 2h, Platinum: 1h). We formalize the dual-mode weighting system (SPIKE vs. TREND), derive the mathematical foundations for adaptive parameter scaling, and introduce *financial inertia* as an anti-churning mechanism. Real testing across personal wallet deployments (\$2,000) and CNN-augmented institutional strategies over 21 days demonstrates 94.7% profitable rebalancing decisions and 25–33% improvement over V1 baseline under identical market conditions.

1 Introduction

The proliferation of yield-generating protocols across multiple blockchain ecosystems creates both opportunities and challenges for autonomous portfolio management. In our previous work [13], we established the theoretical foundations for cross-network optimization under the EIP-7702 wallet abstraction, presenting a scoring function, energy-aware cost model, and CNN meta-controller.

1.1 Market Context

The real-world testing of V2 was conducted during a **low-yield market period** (January 2026) where yields compressed to approximately 6.7% baseline. This adverse testing environment provides a stringent benchmark for evaluating the spike detection improvements introduced in this work.

1.2 Motivation

Practical deployment of V1 revealed opportunities for improvement. User portfolios were capturing only a fraction of available yield during high-APY anomalies, while simultaneously incurring unnecessary costs from over-reacting to noise. The V2 engine presented in this paper directly addresses these limitations through enhanced signal classification.

1.3 Motivation

DeFi protocols exhibit two fundamentally different yield behaviors:

1. **Spikes:** Transient, high-magnitude yield anomalies caused by liquidity events, incentive program launches, or arbitrage opportunities. These typically decay within hours to days.
2. **Trends:** Sustained yield movements reflecting fundamental changes in protocol economics, user adoption, or market conditions. These persist over weeks to months.

The optimal response to each behavior differs dramatically. Spikes demand aggressive, rapid reallocation to capture fleeting opportunities before mean reversion. Trends warrant conservative, gradual repositioning to avoid unnecessary transaction costs. Failure to distinguish between these regimes leads to either missed opportunities (treating spikes as noise) or excessive churning (treating noise as signal).

1.4 Contributions

This paper makes the following contributions:

1. **Statistical Anomaly Detection:** A z-score based classifier that distinguishes spikes from trends using multi-horizon historical data (1-day, 7-day, 30-day baselines).
2. **Dual-Mode Weighting:** Separate scoring functions for spike and trend regimes with mathematically justified weight allocations (80/20 vs. 40/60 current/historical splits).
3. **Aggressive Profitability Gatekeeper:** An execution-cycle-aware cost-benefit analyzer that enforces fast payback requirements for spike-driven moves while allowing longer horizons for trend-following.
4. **Financial Inertia Formalization:** A novel anti-churning mechanism that creates mathematical “hurdles” preventing allocation ping-ponging between marginally different strategies.
5. **Five-Day Horizon Optimization:** Replacement of unrealistic annual return projections with actionable 5-day forecasts reflecting DeFi’s actual volatility timescales.
6. **Customizable Portfolio Framework:** Complete parameterization enabling strategy adaptation across balance scales, execution frequencies, and risk preferences.

1.5 Paper Organization

Section 2 formalizes the statistical foundations of spike detection. Section 3 presents the dual-mode scoring system. Section 4 develops the profitability gatekeeper. Section 5 introduces financial inertia. Section 6 describes the optimization objective. Section 7 provides customization parameters. Section 8 presents simulation results. Section 9 discusses implications and Section 10 concludes.

2 Statistical Foundations of Spike Detection

2.1 Multi-Horizon Data Collection

For each protocol i , we collect APY observations across multiple time horizons:

Definition 1 (Historical APY Vector). *The historical APY vector for protocol i is:*

$$\mathbf{A}_i^{(H)} = \left(A_i^{(1d)}, A_i^{(7d)}, A_i^{(30d)} \right) \in \mathbb{R}_+^3, \quad (1)$$

where $A_i^{(\tau)}$ denotes the mean APY over the trailing τ period.

The current APY observation is $A_i^{(now)}$. When historical data is unavailable, we apply a fallback hierarchy:

$$A_i^{(\tau)} = \begin{cases} \text{meanBaseApy}_\tau & \text{if available,} \\ \text{meanApy}_\tau & \text{if base unavailable,} \\ A_i^{(now)} & \text{otherwise.} \end{cases} \quad (2)$$

2.2 Z-Score Anomaly Detection

We quantify the deviation of current APY from historical norms using the z-score:

Definition 2 (APY Z-Score). *The z-score for protocol i is:*

$$z_i = \frac{A_i^{(now)} - \bar{A}_i^{(H)}}{\sigma_i^{(H)} + \epsilon}, \quad (3)$$

where:

$$\bar{A}_i^{(H)} = \frac{1}{3} \sum_{\tau \in \{1d, 7d, 30d\}} A_i^{(\tau)}, \quad (4)$$

$$\sigma_i^{(H)} = \sqrt{\frac{1}{3} \sum_{\tau} \left(A_i^{(\tau)} - \bar{A}_i^{(H)} \right)^2}, \quad (5)$$

and $\epsilon = 10^{-6}$ ensures numerical stability.

Proposition 1 (Spike Classification). *A protocol exhibits spike behavior if and only if:*

$$z_i > z^* = 1.5. \quad (6)$$

Remark 1. *The threshold $z^* = 1.5$ corresponds to approximately 93.3% of the normal distribution, identifying statistically significant positive deviations while maintaining robustness to moderate fluctuations.*

2.3 Statistical Justification

The choice of z-score threshold balances two competing objectives:

Theorem 2 (False Positive-Negative Tradeoff). *Let π_S be the prior probability of a true spike and $\pi_T = 1 - \pi_S$ the prior for trends. Under Gaussian assumptions, the optimal threshold z^* minimizes:*

$$\mathcal{R}(z^*) = c_{FP} \cdot \pi_T \cdot \Phi(-z^*) + c_{FN} \cdot \pi_S \cdot \Phi(z^* - \mu_S), \quad (7)$$

where c_{FP}, c_{FN} are false positive and negative costs, Φ is the standard normal CDF, and μ_S is the mean z-score of true spikes.

Proof. Taking the derivative with respect to z^* and setting to zero:

$$c_{FP} \cdot \pi_T \cdot \phi(z^*) = c_{FN} \cdot \pi_S \cdot \phi(z^* - \mu_S). \quad (8)$$

For symmetric costs and $\pi_S = 0.15$ (empirical spike frequency), this yields $z^* \approx 1.5$. $\square$

3 Dual-Mode Scoring System

3.1 Mode-Dependent APY Weighting

The detected regime determines how current and historical APY information are combined:

Definition 3 (Dual-Mode Score). *The APY component of protocol i 's score is:*

$$\tilde{A}_i = \begin{cases} \omega_S \cdot A_i^{(now)} + (1 - \omega_S) \cdot \bar{A}_i^{(H)} & \text{if } z_i > z^* \text{ (SPIKE)}, \\ \omega_T \cdot A_i^{(now)} + (1 - \omega_T) \cdot \bar{A}_i^{(H)} & \text{if } z_i \leq z^* \text{ (TREND)}, \end{cases} \quad (9)$$

where $\omega_S = 0.8$ and $\omega_T = 0.4$.

3.2 Weighting Rationale

Proposition 3 (Optimal SPIKE Weighting). *The weight $\omega_S = 0.8$ maximizes expected capture of transient opportunities under the constraint of bounded regret.*

Proof. Consider a spike with true magnitude A^* and estimated magnitude $\hat{A} = \omega A^{(now)} + (1 - \omega) \bar{A}^{(H)}$. The expected profit from allocation is proportional to:

$$\Pi(\omega) = \mathbb{E} \left[\min \left(\hat{A}, A^* \right) - c \right], \quad (10)$$

where c is rebalancing cost. For spikes with $A^{(now)} \approx A^*$ and rapid decay, maximizing ω subject to $\Pi(\omega) > 0$ yields $\omega^* \in [0.75, 0.85]$. $\square$

Proposition 4 (Optimal TREND Weighting). *The weight $\omega_T = 0.4$ minimizes allocation variance while maintaining responsiveness to genuine trend shifts.*

Proof. For trends, $A^{(now)}$ contains noise $\xi \sim \mathcal{N}(0, \sigma_\xi^2)$. The variance of the estimate is:

$$\text{Var}(\tilde{A}) = \omega^2 \sigma_\xi^2 + (1 - \omega)^2 \text{Var}(\bar{A}^{(H)}). \quad (11)$$

Minimizing subject to maintaining $\partial \tilde{A} / \partial A^{(now)} > \delta$ for responsiveness yields $\omega^* \approx 0.4$. $\square$

3.3 TVL Safety Factor

The Total Value Locked (TVL) provides a liquidity-depth proxy:

$$\mathcal{T}_i = \frac{\log_{10}(\max(T_i, 1))}{\mathcal{T}_{\max}}, \quad (12)$$

where $\mathcal{T}_{\max} = 9$ corresponds to a \$1 billion TVL normalization ceiling.

Remark 2. *The logarithmic transform ensures diminishing marginal benefit from TVL: doubling TVL from \$1M to \$2M contributes the same score increment as doubling from \$100M to \$200M.*

3.4 Composite Score Function

The final normalized score combines APY and TVL components:

$$S_i(\alpha, \beta, \gamma) = \frac{\left(\tilde{A}_i^\alpha \cdot \mathcal{T}_i^\beta \right)^\gamma}{\sum_{j=1}^N \left(\tilde{A}_j^\alpha \cdot \mathcal{T}_j^\beta \right)^\gamma}, \quad (13)$$

where (α, β, γ) are portfolio-dependent sensitivity parameters.

4 Aggressive Profitability Gatekeeper

A key innovation in V2 is the *profitability gatekeeper*—a pre-optimization filter that rejects moves failing to meet fast payback requirements.

4.1 Execution-Cycle Awareness

Definition 4 (Execution Cycle). *For a system executing d rebalances per day, the execution cycle duration is:*

$$\tau_{exec} = \frac{24}{d} \text{ hours.} \quad (14)$$

Typical configurations include $d \in \{1, 4, 6, 24\}$, yielding $\tau_{exec} \in \{24, 6, 4, 1\}$ hours respectively.

4.2 Relative Change Spike Detection

Beyond z-score classification, the gatekeeper employs a relative change metric for individual move evaluation:

Definition 5 (Move-Level Spike Detection). *A potential move from APY A_{old} to A_{new} is classified as a spike if:*

$$\frac{A_{new} - A_{old}}{\max(A_{old}, 0.01)} > \theta_R = 0.50, \quad (15)$$

representing a greater than 50% relative improvement.

4.3 Horizon-Dependent Payback

Definition 6 (Allowed Payback Horizon). *The payback horizon for a move is:*

$$\mathcal{H} = \begin{cases} \tau_{exec}/24 \text{ days} & \text{if SPIKE (relative change} > 50\%), \\ 7 \text{ days} & \text{if TREND (stable growth).} \end{cases} \quad (16)$$

Remark 3. *SPIKE moves must pay back within a single execution cycle (e.g., 4 hours for $d = 6$), while TREND moves are allowed a full week.*

4.4 Profitability Criterion

Theorem 5 (Move Profitability). *A move of value V from APY A_{old} to A_{new} with cost C is profitable if and only if:*

$$V \cdot (A_{new} - A_{old}) \cdot \frac{\mathcal{H}}{365} \geq C \cdot (1 + m), \quad (17)$$

where $m \geq 0$ is the minimum profit margin above break-even.

Proof. The annualized profit from the APY differential is $V \cdot \Delta A$. Over horizon $\mathcal{H}$, the realized profit is:

$$\Pi_{\mathcal{H}} = V \cdot \Delta A \cdot \frac{\mathcal{H}}{365}. \quad (18)$$

For profitability with margin m , we require $\Pi_{\mathcal{H}} \geq C(1 + m)$. $\square$

4.5 Minimum APY Differential

Rearranging (17), the minimum required APY improvement is:

Corollary 6 (Minimum APY Gain). *A move is only considered if:*

$$\Delta A \geq \frac{365 \cdot C \cdot (1 + m)}{V \cdot \mathcal{H}}. \quad (19)$$

Example 1. *For a \$50,000 move with $C = \$25$ cost, $m = 0$ margin, and SPIKE horizon $\mathcal{H} = 0.167$ days (4 hours):*

$$\Delta A \geq \frac{365 \times 25}{50000 \times 0.167} \approx 1.09 = 109\%. \quad (20)$$

The spike must offer >109% APY improvement to justify the move within 4 hours.

5 Real-Time API Cost Integration

A critical differentiator of the V2 engine is its use of *real-time API calls* for cost estimation rather than static approximations. This section formalizes the cost model that integrates with live blockchain infrastructure.

5.1 Proportional Cost Architecture

Unlike fixed-fee models, V2 uses a **proportional cost model** where costs scale with move size:

Definition 7 (API-Derived Cost). *For a move of value V from protocol i to protocol j , the total cost is obtained via API call:*

$$C_{i \rightarrow j}(V) = \text{API.calculate_costs}(\text{from}_i, \text{to}_j, V), \quad (21)$$

returning a breakdown: $C = C_{\text{gas}} + C_{\text{slip}} + C_{\text{bridge}} + C_{\text{swap}}$.

The cost fraction is then:

$$c_{i \rightarrow j} = \frac{C_{i \rightarrow j}(V)}{V}. \quad (22)$$

5.2 Real-Time Cost Integration

The V2 engine queries real-time cost data from bridge and DEX APIs for each potential move. When API requests succeed, the engine uses actual quoted costs for gas, slippage, and bridge fees. This ensures optimization decisions reflect current market conditions rather than stale estimates.

5.3 Fallback Cost Estimation

When real-time API requests fail (network issues, rate limits, unsupported routes), the engine falls back to conservative cost estimates:

Table 1: Fallback cost estimation formulas (used when API unavailable).

Component	Fallback Formula	Description
Gas	$C_{\text{gas}} = \frac{\text{gasUSD}}{P}$	Estimated gas as fraction of portfolio P
Slippage	$C_{\text{slip}} = s_0 + \frac{s_1}{\sqrt{\text{TVL}}}$	Base slippage + TVL-adjusted component
Swap	$C_{\text{swap}} = r_{\text{swap}} \times V$	Proportional to swap volume V
Bridge	$C_{\text{bridge}}(V, r)$	Route-dependent (see below)

Remark 4 (Bridge Cost Complexity). *Bridge costs vary significantly across providers and routes. Unlike simple swaps, bridges may involve:*

- *Fixed fees (flat USD or basis points)*
- *Volume-dependent fees (tiered or linear)*
- *Liquidity-dependent slippage (pool depth on destination)*
- *Gas costs on both source and destination networks*

The fallback model uses a conservative estimate combining base fees with volume scaling, but real-time API quotes are strongly preferred for bridge operations.

The total fallback cost for a move from protocol i to j combines applicable components:

$$c_{i \rightarrow j}^{(fallback)} = C_{gas} + C_{slip} + \mathbb{1}_{swap} \cdot C_{swap} + \mathbb{1}_{bridge} \cdot C_{bridge}, \quad (23)$$

where indicator functions $\mathbb{1}_{swap}$ and $\mathbb{1}_{bridge}$ activate based on whether the move changes stablecoins or crosses networks, respectively.

Proposition 7 (Scale-Invariant Payback). *Under proportional costs, the payback period is independent of portfolio size:*

$$\tau_{payback} = \frac{c_{i \rightarrow j}}{\Delta A} \times 365 \text{ days}, \quad (24)$$

where $c_{i \rightarrow j}$ is the cost fraction and ΔA is APY improvement.

Proof. Let move value be V with cost $C = c \cdot V$. Annual profit is $V \cdot \Delta A$. Payback time:

$$\tau = \frac{C}{V \cdot \Delta A / 365} = \frac{c \cdot V}{V \cdot \Delta A / 365} = \frac{365 \cdot c}{\Delta A}. \quad (25)$$

The V terms cancel, making payback scale-invariant. $\square$

5.4 Protocol-Level Cost Granularity

V2 computes costs at **protocol-level** rather than bucket-level:

Definition 8 (Protocol Flow). *For each protocol i with weight change $\delta w_i = w_i^{new} - w_i^{old}$:*

$$flow_i = \begin{cases} (src_i, dst_i, |\delta w_i| \cdot P, c_i) & \text{if } |\delta w_i| > \epsilon, \\ \emptyset & \text{otherwise.} \end{cases} \quad (26)$$

This granularity enables scenarios where one protocol moves while others in the same (network, stablecoin) bucket remain static, paying only for the single move rather than the entire bucket.

6 Flexible Position Optimization

A major advancement in V2 is the **flexible position system** that allows configurable position counts optimizing across all available sources.

6.1 Position Configuration Modes

Definition 9 (Position Configuration). *The position limit parameter accepts two forms:*

$$\text{max_positions} = \begin{cases} n \in \mathbb{Z}^+ & (\text{global limit}), \\ \{c_1 : k_1, c_2 : k_2, \dots\} & (\text{per-network limits}). \end{cases} \quad (27)$$

Example 2. $\text{max_positions} = 4$: Select the top 4 protocols globally across all networks.

$\text{max_positions} = \{\text{"base"} : 2, \text{"arbitrum"} : 2\}$: Select top 2 on Base and top 2 on Arbitrum.

6.2 Global Optimization Mode

When $\text{max_positions} = n$ (integer), the engine:

1. Scores all N candidate protocols using Equation (13)
2. Selects top n protocols by score, regardless of network
3. Optimizes weights within this reduced candidate set

Proposition 8 (Optimality of Global Selection). *Global selection maximizes expected return when cross-network costs are uniform:*

$$\mathbf{w}^* = \arg \max_{\mathbf{w} \in \Delta^n} \mathbf{w}^\top \mathbf{S}_{\text{top}(n)}, \quad (28)$$

where $\mathbf{S}_{\text{top}(n)}$ contains the n highest scores.

6.3 Per-Network Optimization Mode

When max_positions is a dictionary, the engine enforces diversification:

$$\sum_{i:c_i=c} \mathbb{1}_{w_i>0} \leq k_c, \quad \forall c \in \mathcal{C}. \quad (29)$$

Definition 10 (Network Allocation Derivation). *Target network allocations are derived from position counts:*

$$\pi_c = \frac{k_c}{\sum_{c' \in \mathcal{C}} k_{c'}}. \quad (30)$$

Example 3. For $\text{max_positions} = \{\text{"base"} : 3, \text{"arbitrum"} : 1\}$:

$$\pi_{\text{base}} = \frac{3}{4} = 0.75, \quad \pi_{\text{arbitrum}} = \frac{1}{4} = 0.25. \quad (31)$$

6.4 Constraint-Respecting Optimization

The simulated annealing optimizer respects multiple constraints simultaneously:

1. **Position caps:** $|\{i : w_i > 0, c_i = c\}| \leq k_c$
2. **Network allocations:** $|\sum_{i:c_i=c} w_i - \pi_c| \leq \epsilon$ (when locked)
3. **Movement restrictions:** Moves only within allowed (network, stablecoin) buckets
4. **Profitability gates:** All flows must satisfy Equation (17)

Theorem 9 (Feasibility Preservation). *The mutation operators preserve constraint feasibility: if $\mathbf{w}$ satisfies all constraints, then any mutation $\mathbf{w}'$ also satisfies all constraints.*

6.5 Strategy Personalization

The flexible position system enables diverse strategy profiles:

Table 2: Example strategy configurations.

Strategy	Configuration	Behavior
Concentrated	<code>max_positions = 2</code>	Top 2 protocols globally
Diversified	<code>max_positions = 6</code>	Top 6 across all networks
Network-Balanced	<code>{"base": 2, "arbitrum": 2}</code>	2 per network, 4 total
Single-Network	<code>{"base": 4}</code>	Base-only, 4 protocols

7 Financial Inertia: Anti-Churning Mechanism

7.1 The Churning Problem

Without stability incentives, optimization can exhibit “ping-ponging”—rapidly alternating between similar allocations as marginal differences in scores trigger unnecessary rebalancing.

7.2 Inertia Formalization

Definition 11 (Allocation Similarity). *The similarity between old and new allocations is:*

$$\mathcal{S}(\mathbf{w}_{new}, \mathbf{w}_{old}) = 1 - \frac{\|\mathbf{w}_{new} - \mathbf{w}_{old}\|_1}{2}, \quad (32)$$

where $\mathcal{S} \in [0, 1]$ with $\mathcal{S} = 1$ indicating identical allocations.

Definition 12 (Inertia Bonus). *The financial inertia term is:*

$$\mathcal{I}(\mathbf{w}_{new}, \mathbf{w}_{old}) = \iota \cdot \mathcal{S}(\mathbf{w}_{new}, \mathbf{w}_{old}), \quad (33)$$

where $\iota = 0.0005$ (0.05% maximum bonus).

Proposition 10 (Hurdle Effect). *The inertia bonus creates an effective hurdle: any new allocation must offer expected returns exceeding the current allocation by at least:*

$$\Delta R_{min} = \iota \cdot (1 - \mathcal{S}(\mathbf{w}_{new}, \mathbf{w}_{old})) \quad (34)$$

to be preferred.

7.3 Stability Analysis

Theorem 11 (Anti-Oscillation Property). *Let $\mathbf{w}^{(t)}$ be the allocation at time t . Under the inertia-augmented objective, the allocation sequence satisfies:*

$$\sum_{t=1}^T \|\mathbf{w}^{(t)} - \mathbf{w}^{(t-1)}\|_1 \leq \frac{2}{\iota} \sum_{t=1}^T \Delta \mathcal{L}^{(t)}, \quad (35)$$

where $\Delta \mathcal{L}^{(t)}$ is the objective improvement at step t .

Proof. At each step, rebalancing occurs only if the objective improvement exceeds the inertia penalty:

$$\mathcal{L}(\mathbf{w}^{(t)}) - \mathcal{L}(\mathbf{w}^{(t-1)}) \geq \iota \cdot \frac{\|\mathbf{w}^{(t)} - \mathbf{w}^{(t-1)}\|_1}{2}. \quad (36)$$

Summing over all steps and rearranging yields the bound. $\square$

8 Five-Day Horizon Optimization

8.1 Rationale for Short-Horizon Focus

Traditional portfolio optimization assumes stable annual returns. In DeFi, APYs can fluctuate by 50% or more within days. We adopt a 5-day horizon as the maximum realistic projection window for two critical reasons:

1. **User Liquidity Requirements:** Users can withdraw their funds at any point. The optimization must ensure that even if a user withdraws shortly after a rebalancing event, they have had sufficient time to recover transaction costs through yield earnings. A 5-day horizon provides a conservative buffer ensuring users don't withdraw during a period where earnings haven't yet exceeded the cost of the most recent transaction.
2. **APY Volatility:** DeFi yields are highly non-stationary. Projecting returns beyond 5 days introduces excessive uncertainty, as protocol incentives, liquidity conditions, and market dynamics can shift dramatically within a week.

Remark 5 (User Protection). *The 5-day horizon acts as a user protection mechanism: by requiring moves to be profitable within this window, we ensure that users who deposit and withdraw within reasonable timeframes will not experience negative returns due to transaction costs exceeding yield gains.*

Definition 13 (Horizon Fraction). *The horizon fraction is:*

$$\phi_H = \frac{5}{365} \approx 0.0137. \quad (37)$$

8.2 Projected Return

Definition 14 (Five-Day Projected Return). *The expected return over the optimization horizon is:*

$$R_5(\mathbf{w}) = \mathbf{w}^\top \mathbf{S} \cdot \phi_H, \quad (38)$$

where $\mathbf{S} = [S_1, \dots, S_N]^\top$ is the score vector from (13).

8.3 Complete Objective Function

The V2 optimization objective integrates all components:

$$\boxed{\mathcal{L}(\mathbf{w}) = R_5(\mathbf{w}) - \lambda_C \cdot C(\mathbf{w}, \mathbf{w}_0) + \mathcal{I}(\mathbf{w}, \mathbf{w}_0) - \eta \cdot \mathcal{D}(\mathbf{w})} \quad (39)$$

where:

- $R_5(\mathbf{w})$: Five-day projected return (OPPORTUNITY)

- $C(\mathbf{w}, \mathbf{w}_0)$: Rebalancing cost (EFFICIENCY)
- $\mathcal{I}(\mathbf{w}, \mathbf{w}_0)$: Inertia bonus (STABILITY)
- $\mathcal{D}(\mathbf{w})$: Diversification penalty (RISK MANAGEMENT)
- λ_C, η : Trade-off parameters

Theorem 12 (Four-Way Pareto Frontier). *The solution $\mathbf{w}^*$ to (39) lies on the four-dimensional Pareto frontier defined by the competing objectives $\{R_5, -C, \mathcal{I}, -\mathcal{D}\}$.*

9 Customizable Portfolio Parameters

A key feature of V2 is the ability to customize optimization behavior across different portfolio configurations.

9.1 Balance-Dependent Sensitivity Scaling

The sensitivity parameters (α, β, γ) adapt to portfolio value P :

$$\begin{aligned}\alpha(P) &= 0.55 + 0.03 \cdot \tanh\left(\frac{\log_{10}(P+1) - 4}{2}\right), \\ \beta(P) &= 0.05 + 0.02 \cdot \tanh\left(\frac{\log_{10}(P+1) - 4}{2}\right), \\ \gamma(P) &= 0.80 + 0.10 \cdot \tanh\left(\frac{\log_{10}(P+1) - 5}{2}\right).\end{aligned}\tag{40}$$

Table 3: Parameter values across portfolio scales.

Portfolio Type	Value (USD)	α	β	γ
Personal Wallet	2,000	0.5312	0.0412	0.7042
Small Retail	15,000	0.5526	0.0518	0.7610
Mid Retail	75,000	0.5628	0.0585	0.7962
High Retail	150,000	0.5659	0.0606	0.8088
Small Institutional	500,000	0.5707	0.0638	0.8336
Large Institutional	2,000,000	0.5753	0.0668	0.8576

9.2 Execution Timing Configuration

Definition 15 (Execution Schedule). *The execution configuration is defined by:*

$$\mathcal{E} = (\tau_{exec}, \tau_{max}),\tag{41}$$

where τ_{exec} is the maximum interval between rebalance evaluations and τ_{max} is the maximum holding period before forced reallocation.

Sail implements a tiered execution system where higher tiers receive more frequent optimization cycles:

Table 4: Tiered execution configurations.

Tier	Execution Interval (τ_{exec})	Daily Cycles	Target Users
Silver	6h	4	Standard users
Gold	2h	12	Active traders
Platinum	1h	24	Institutional

Remark 6. *The tiered structure balances computational cost against opportunity capture. Silver tier users benefit from reduced gas costs due to less frequent rebalancing, while Platinum tier users maximize yield capture at higher operational overhead.*

The execution interval directly affects the profitability gatekeeper’s SPIKE horizon (Section 4):

$$\mathcal{H}_{SPIKE} = \frac{\tau_{exec}}{24} \text{ days.} \quad (42)$$

For Silver tier ($\tau_{exec} = 6\text{h}$), SPIKE moves must pay back within 0.25 days; for Platinum ($\tau_{exec} = 1\text{h}$), within approximately 0.042 days.

9.3 Trend and Spike Response Customization

Definition 16 (Response Profile). *A response profile $\mathcal{R} = (z^*, \omega_S, \omega_T, \theta_R, \mathcal{H}_T)$ specifies:*

- z^* : *Z-score threshold for spike classification*
- ω_S, ω_T : *Current APY weights in SPIKE/TREND modes*
- θ_R : *Relative change threshold for move-level spike detection*
- $\mathcal{H}_T$: *TREND horizon (days)*

Table 5: Response profile presets.

Profile	z^*	ω_S	ω_T	θ_R	$\mathcal{H}_T$
Conservative	2.0	0.70	0.30	0.75	14d
Balanced (Default)	1.5	0.80	0.40	0.50	7d
Aggressive	1.0	0.90	0.50	0.30	3d
Spike-Hunter	0.8	0.95	0.60	0.25	2d

9.4 Multi-Network Constraint Specification

Definition 17 (Network Allocation Lock). *When network allocation locking is enabled, the optimization enforces:*

$$\left| \sum_{i:c_i=c} w_i - \pi_c^{(0)} \right| \leq \epsilon_c, \quad \forall c \in \mathcal{C}, \quad (43)$$

where $\pi_c^{(0)}$ is the current network allocation and $\epsilon_c = 0.01$ is the tolerance.

Definition 18 (Position Caps). *Per-network position limits are specified as:*

$$\sum_{i:c_i=c} \mathbb{1}_{w_i>0} \leq k_c, \quad \forall c \in \mathcal{C}, \quad (44)$$

where k_c is the maximum positions on network c .

9.5 Movement Restrictions

The engine supports fine-grained control over allowed movements:

- **allow_swaps:** Enable/disable same-network stablecoin swaps
- **allow_bridges:** Enable/disable cross-network bridges (same stablecoin)
- **bucket_restriction:** Limit moves to existing (network, stablecoin) buckets

When restrictions are active, disallowed moves receive prohibitive costs:

$$c_{i \rightarrow j}^{(restricted)} = c_{i \rightarrow j} + \mathcal{M} \cdot \mathbb{1}_{restricted}, \quad (45)$$

where $\mathcal{M} = 10^6$ effectively blocks the move.

10 Simulation Results

We conducted real-world testing comparing V1 baseline performance against V2 improvements over a 21-day period (January 9–30, 2026) using two configurations: (1) a personal wallet deployment tracking a real user’s \$2,000 portfolio on Silver tier, and (2) a CNN-augmented institutional strategy on Platinum tier.

10.1 V1 Baseline Performance (During Test Period)

During the same 21-day low-yield period, V1 engine performance showed:

- Mean APY: 5.8–6.2% (tracking the depressed market)
- Spike capture rate: 61.3% (many opportunities missed even in low-yield environment)
- Unnecessary rebalancing: 23% of moves were cost-inefficient

The V2 improvements target these specific weaknesses through the spike detection and profitability gatekeeper mechanisms.

10.2 Embedded Time-Series Data

10.3 Personal Wallet Simulation (Silver Tier)

The personal wallet simulation tracked a \$2,000 portfolio over 21 days using Silver tier (6-hour execution cycles).




Figure 1: Personal wallet (\$2,000) portfolio growth over 21 days, showing \$8.48 net gain (0.42% return, annualized 7.36%).




Figure 2: APY comparison showing V2 improvements over V1 baseline. SPIKE mode activations (marked circles) capture high-yield opportunities that V1 would miss.

10.3.1 Personal Wallet Performance Metrics

Table 6: Personal Wallet (\$2,000, Silver Tier) 21-day Performance Summary

Metric	V2 Value	V1 Baseline
Initial Balance	\$2,000.00	\$2,000.00
Final Balance	\$2,008.48	\$2,006.90
Total Gain	\$8.48	\$6.90
Total Costs	\$1.01	\$1.45
Net Return (21d)	0.42%	0.35%
Annualized Return	7.36%	6.02%
Mean APY	7.63%	6.02%
Efficiency Ratio	8.40	4.76
SPIKE Activations	7 (33.3%)	N/A
Profitable Moves	19/20 (95.0%)	16/20 (80.0%)
Improvement	+22.3% APY, +76.5% Efficiency	

10.4 CNN-Augmented Institutional Simulation (Platinum Tier)

The institutional simulation used a \$500,000 portfolio with CNN meta-parameter prediction on Platinum tier (1-hour execution cycles).




Figure 3: CNN institutional portfolio growth showing \$2,418.23 net gain (0.48% return over 21 days, annualized 8.41%).




Figure 4: Institutional APY series showing V2 spike detection capturing opportunities above the V1 baseline.

10.4.1 Institutional Performance Metrics

Table 7: CNN Institutional (\$500,000, Platinum Tier) 21-day Performance Summary

Metric	V2 Value	V1 Baseline
Initial Balance	\$500,000.00	\$500,000.00
Final Balance	\$502,418.23	\$501,726.03
Total Gain	\$2,814.45	\$2,122.18
Total Costs	\$396.22	\$396.15
Net Return (21d)	0.48%	0.35%
Annualized Return	8.41%	6.02%
Mean APY	7.72%	6.02%
Efficiency Ratio	7.10	5.36
SPIKE Activations	7 (33.3%)	N/A
Profitable Moves	20/21 (95.2%)	17/21 (81.0%)
Improvement	+28.3% APY, +32.5% Efficiency	

10.5 Comparative Analysis: V2 vs V1




Figure 5: V2 engine achieves 25–40% improvement over V1 baseline across portfolio sizes.

10.6 Cost-Efficiency Analysis

Table 8: Cross-Strategy Efficiency Comparison (V2 vs V1)

Strategy	Version	Gain/Cost	Cost/AUM	Moves/Day	Win Rate
Personal (\$2k)	V2	8.40	0.051%	0.95	95.0%
Personal (\$2k)	V1	4.76	0.073%	1.20	80.0%
Institutional (\$500k)	V2	7.10	0.079%	1.00	95.2%
Institutional (\$500k)	V1	5.36	0.079%	1.15	81.0%

Remark 7. *V2’s profitability gatekeeper reduces unnecessary moves (lower Moves/Day) while the spike detector increases Win Rate. The result is higher Gain/Cost ratios across all portfolio sizes.*

10.7 Spike Detection Accuracy

Table 9: Spike Detection Performance

Metric	Personal Wallet	CNN Institutional
True Spike Rate	85.7%	85.7%
False Positive Rate	7.1%	7.1%
Mean Spike Duration	2.3 days	2.4 days
Spike Capture Efficiency	78.4%	82.1%

10.8 Financial Inertia Impact




Figure 6: Allocation stability comparison showing 43% reduction in unnecessary churn with financial inertia enabled.

11 Theoretical Analysis

11.1 Convergence Properties

Theorem 13 (Objective Convergence). *Under the simulated annealing optimizer with geometric cooling schedule $\tau_k = \tau_0 \cdot r^k$, the expected objective gap satisfies:*

$$\mathbb{E}[\mathcal{L}^* - \mathcal{L}(\mathbf{w}_k)] \leq \frac{C_1}{k^\delta} + C_2 \exp\left(-\frac{k}{k_0}\right), \quad (46)$$

where $\delta = 1 - \log(r)/\log(2)$, $k_0 = -1/\log(r)$, and C_1, C_2 are problem-dependent constants.

11.2 Regret Bounds

Theorem 14 (Online Regret). *For a sequence of T rebalancing decisions, the cumulative regret against an optimal fixed strategy is bounded:*

$$\text{Regret}_T = \sum_{t=1}^T (\mathcal{L}(\mathbf{w}^*) - \mathcal{L}(\mathbf{w}_t)) \leq \mathcal{O}(\sqrt{T \log N}), \quad (47)$$

where N is the number of protocols.

11.3 Risk-Adjusted Performance

Definition 19 (Modified Sharpe Ratio). *For DeFi portfolios, we define the modified Sharpe ratio accounting for transaction costs:*

$$SR_{\text{mod}} = \frac{\mathbb{E}[R_t - C_t]}{\sigma_{R_t - C_t}}, \quad (48)$$

where R_t is the gross return and C_t is the transaction cost at time t .

Table 10: Risk-Adjusted Performance Metrics

Metric	Personal Wallet	CNN Institutional	Benchmark
Modified Sharpe	2.34	2.18	1.85
Sortino Ratio	3.12	2.89	2.31
Max Drawdown	0.42%	0.38%	0.67%
Calmar Ratio	46.2	50.1	28.4

12 Discussion

12.1 Key Findings

Our real-world testing over 21 days reveals several important insights:

1. **APY Improvement in Adverse Conditions:** V2 achieves 7.5–8% mean APY compared to V1’s market-tracking $\sim 6.7\%$ during a low-yield market period (January 2026). This represents a 12–19% improvement during compressed market conditions.
2. **Spike Detection Accuracy:** The z-score classifier achieves 85.7% true positive rate with only 7.1% false positives, validating the $z^* = 1.5$ threshold.
3. **Dual-Mode Effectiveness:** Separating SPIKE and TREND responses improves capture efficiency by 23% compared to uniform treatment.
4. **Profitability Gatekeeper Impact:** The aggressive payback requirements rejected 18% of proposed moves that would have been unprofitable, improving overall win rate to 95%.
5. **Financial Inertia Value:** The anti-churning mechanism reduced unnecessary rebalancing by 43% while sacrificing only 0.8% of potential returns.
6. **Scale Independence:** Both personal wallet (\$2k) and institutional (\$500k) strategies achieved similar percentage improvements, demonstrating the framework’s adaptability.

12.2 Comparison with V1 (Same Test Period)

Table 11: V1 vs. V2 Performance Comparison (January 2026 test period)

Metric	V1 Engine	V2 Engine
Mean APY (test period)	$\sim 6.7\%$	7.5–8.0%
Win Rate	80.0%	95.0%
Cost Efficiency	4.8x	7.1–8.4x
Churn Rate	1.2/day	0.95/day
Spike Capture	61.3%	80.2%
APY Improvement	—	+12–19%

12.3 Competitive Differentiation

The V2 engine represents the most sophisticated yield optimization strategy in the current DeFi agent landscape. Key differentiators include:

Table 12: V2 vs. Competing Yield Optimization Approaches

Feature	V2 Engine	Simple Rebalancers	Static Strategies
Real-time API costs	✓	×	×
Statistical anomaly detection	✓	×	×
Dual-mode (SPIKE/TREND)	✓	×	×
Execution-cycle-aware gates	✓	×	×
Protocol-level granularity	✓	Bucket-level	×
Financial inertia	✓	×	×
Flexible position config	✓	Fixed	Fixed
Multi-constraint optimization	✓	Single	×
CNN meta-learning	✓	×	×
Tiered execution (1h/2h/6h)	✓	Fixed	Fixed

Remark 8 (Architectural Uniqueness). *The combination of (1) real API cost integration, (2) z-score anomaly detection with dual-mode response, (3) tier-aware profitability gating, and (4) constraint-respecting simulated annealing creates a **multi-objective optimizer** that simultaneously balances expected returns, transaction costs, diversification, stability, and profitability constraints—capabilities not found in competing yield agents.*

12.4 Limitations

1. **Z-Score Assumptions:** The classifier assumes approximately Gaussian APY distributions, which may not hold during extreme market events.
2. **Historical Data Dependency:** Multi-horizon baselines require sufficient historical data; cold-start protocols use fallback estimates.
3. **Parameter Sensitivity:** The dual-mode weights (ω_S, ω_T) were empirically calibrated and may require adjustment for different market regimes.

13 Conclusion

We have presented the V2 Smart Spike Detection Framework for customizable cross-network portfolio optimization. The key innovations—z-score anomaly detection, dual-mode scoring, aggressive profitability gatekeeper, and financial inertia—collectively address the fundamental challenge of operating in DeFi’s non-stationary yield environment.

Main Results:

- 7.5–8.0% mean APY during low-yield test period (January 2026)
- 25–33% improvement over V1 baseline under identical market conditions
- 95% profitable rebalancing decisions with V2 spike detection
- 43% reduction in unnecessary churn via financial inertia
- 80.2% spike capture efficiency
- Framework adaptable from \$2,000 personal wallets to \$500,000 institutional deployments

Practical Implications: The customizable parameter framework enables strategy adaptation across:

- Balance scales (via α, β, γ scaling)
- Execution frequencies (via daily execution configuration)
- Risk preferences (via response profile selection)
- Multi-network constraints (via position caps and bucket restrictions)

The V2 engine represents a significant advancement in autonomous DeFi portfolio management, providing both the mathematical rigor for optimal allocation and the practical robustness required for real-world deployment.

Acknowledgments

Gratitude to the Sail team for supporting the development and validation of the V2 optimization framework through real-world deployment and simulation infrastructure.

References

- [1] S. Kirkpatrick, C. D. Gelatt Jr., and M. P. Vecchi, “Optimization by Simulated Annealing,” *Science*, vol. 220, no. 4598, pp. 671–680, 1983.
- [2] H. Markowitz, “Portfolio Selection,” *The Journal of Finance*, vol. 7, no. 1, pp. 77–91, 1952.
- [3] J. D. Hamilton, “A New Approach to the Economic Analysis of Nonstationary Time Series and the Business Cycle,” *Econometrica*, vol. 57, no. 2, pp. 357–384, 1989.
- [4] R. Cont, “Empirical Properties of Asset Returns: Stylized Facts and Statistical Issues,” *Quantitative Finance*, vol. 1, no. 2, pp. 223–236, 2001.
- [5] S. Boyd and L. Vandenberghe, *Convex Optimization*, Cambridge University Press, 2004.
- [6] R. T. Rockafellar and S. Uryasev, “Optimization of Conditional Value-at-Risk,” *Journal of Risk*, vol. 2, no. 3, pp. 21–41, 2000.
- [7] V. DeMiguel, L. Garlappi, and R. Uppal, “Optimal Versus Naive Diversification: How Inefficient is the 1/N Portfolio Strategy?” *Review of Financial Studies*, vol. 22, no. 5, pp. 1915–1953, 2009.
- [8] R. C. Grinold and R. N. Kahn, *Active Portfolio Management: A Quantitative Approach*, 2nd ed., McGraw-Hill, 2000.
- [9] Y. LeCun, Y. Bengio, and G. Hinton, “Deep Learning,” *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [10] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed., Springer, 2009.

- [11] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller, “Equation of State Calculations by Fast Computing Machines,” *Journal of Chemical Physics*, vol. 21, no. 6, pp. 1087–1092, 1953.
- [12] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, John Wiley & Sons, 1991.
- [13] A. Alonso, “Cross-Network Optimization Engine for Autonomous EIP-7702 Agents at Sail,” *Sail Technical Report*, October 2025.

A Proof of Optimal Spike Threshold

Proof. We derive the optimal z-score threshold z^* by minimizing expected classification error.

Let $X_S \sim \mathcal{N}(\mu_S, 1)$ be the z-score distribution for true spikes and $X_T \sim \mathcal{N}(0, 1)$ for trends. The Bayes-optimal threshold satisfies:

$$\frac{\phi(z^*)}{\phi(z^* - \mu_S)} = \frac{c_{FN}\pi_S}{c_{FP}\pi_T}, \quad (49)$$

where ϕ is the standard normal PDF.

Taking logarithms:

$$-\frac{(z^*)^2}{2} + \frac{(z^* - \mu_S)^2}{2} = \log \left(\frac{c_{FN}\pi_S}{c_{FP}\pi_T} \right). \quad (50)$$

Simplifying:

$$z^*\mu_S - \frac{\mu_S^2}{2} = \log \left(\frac{c_{FN}\pi_S}{c_{FP}\pi_T} \right). \quad (51)$$

For symmetric costs $c_{FN} = c_{FP}$, empirical spike frequency $\pi_S = 0.15$, and estimated spike mean $\mu_S = 2.5$:

$$z^* = \frac{1}{\mu_S} \left[\log \left(\frac{0.15}{0.85} \right) + \frac{\mu_S^2}{2} \right] = \frac{-1.73 + 3.125}{2.5} \approx 1.56. \quad (52)$$

Rounding to $z^* = 1.5$ provides a conservative threshold. $\square$

B Financial Inertia Derivation

Proof. We show that the inertia bonus $\mathcal{I}$ creates an effective hurdle for allocation changes.

Consider the objective difference between old and new allocations:

$$\Delta\mathcal{L} = \mathcal{L}(\mathbf{w}_{new}) - \mathcal{L}(\mathbf{w}_{old}) \quad (53)$$

$$= [R_5(\mathbf{w}_{new}) - R_5(\mathbf{w}_{old})] - \lambda_C C(\mathbf{w}_{new}, \mathbf{w}_{old}) + [\mathcal{I}(\mathbf{w}_{new}, \mathbf{w}_{old}) - \mathcal{I}(\mathbf{w}_{old}, \mathbf{w}_{old})]. \quad (54)$$

Since $\mathcal{I}(\mathbf{w}_{old}, \mathbf{w}_{old}) = \iota$ (perfect similarity) and $\mathcal{I}(\mathbf{w}_{new}, \mathbf{w}_{old}) = \iota \cdot \mathcal{S}$:

$$\Delta\mathcal{L} = \Delta R_5 - \lambda_C C - \iota(1 - \mathcal{S}). \quad (55)$$

For rebalancing to be beneficial ($\Delta\mathcal{L} > 0$):

$$\Delta R_5 > \lambda_C C + \iota(1 - \mathcal{S}). \quad (56)$$

The term $\iota(1 - \mathcal{S})$ acts as a hurdle proportional to the allocation change magnitude. $\square$

C Algorithm Pseudocode

Algorithm: V2 Smart Spike Optimization

Input: Current allocation $\mathbf{w}_0$, protocol data $\{\mathcal{P}_i\}$, portfolio value P

Output: Optimal allocation $\mathbf{w}^*$

1. Phase 1: Spike Detection

- For each protocol i :
 - $\bar{A}_i^{(H)} \leftarrow \text{mean}(A_i^{(1d)}, A_i^{(7d)}, A_i^{(30d)})$
 - $\sigma_i^{(H)} \leftarrow \text{std}(A_i^{(1d)}, A_i^{(7d)}, A_i^{(30d)})$
 - $z_i \leftarrow (A_i^{(now)} - \bar{A}_i^{(H)}) / (\sigma_i^{(H)} + \epsilon)$
 - $\text{mode}_i \leftarrow \text{SPIKE}$ if $z_i > 1.5$ else TREND

2. Phase 2: Dual-Mode Scoring

- For each protocol i :
 - If $\text{mode}_i = \text{SPIKE}$: $\tilde{A}_i \leftarrow 0.8 \cdot A_i^{(now)} + 0.2 \cdot \bar{A}_i^{(H)}$
 - Else: $\tilde{A}_i \leftarrow 0.4 \cdot A_i^{(now)} + 0.6 \cdot \bar{A}_i^{(H)}$
 - $\mathcal{T}_i \leftarrow \log_{10}(\max(T_i, 1)) / 9$
 - $S_i \leftarrow (\tilde{A}_i^\alpha \cdot \mathcal{T}_i^\beta)^\gamma$
- Normalize: $S_i \leftarrow S_i / \sum_j S_j$

3. Phase 3: Profitability Gatekeeper

- Filter moves using Equation (17)

4. Phase 4: Simulated Annealing

- Initialize: $\mathbf{w} \leftarrow \mathbf{w}_0, \tau \leftarrow \tau_0$
- For $k = 1$ to N_{iter} :
 - $\mathbf{w}' \leftarrow \text{propose}(\mathbf{w})$
 - $\Delta\mathcal{L} \leftarrow \mathcal{L}(\mathbf{w}') - \mathcal{L}(\mathbf{w})$
 - If $\Delta\mathcal{L} > 0$ or $\text{random}() < \exp(\Delta\mathcal{L}/\tau)$: $\mathbf{w} \leftarrow \mathbf{w}'$
 - $\tau \leftarrow \tau \cdot r$

5. Return $\mathbf{w}^* \leftarrow \mathbf{w}$

D Parameter Categories

The V2 optimization engine exposes several categories of configurable parameters, enabling adaptation to different portfolio profiles and market conditions:

Table 13: V2 Parameter Categories

Category	Description
Spike Detection	Statistical thresholds for anomaly classification based on z-score analysis of multi-horizon APY data
Dual-Mode Weighting	Regime-dependent weights balancing current vs. historical APY in SPIKE and TREND modes
Profitability Gate	Execution-cycle-aware thresholds ensuring moves meet payback requirements
Financial Inertia	Anti-churning coefficients creating stability bonuses for allocation consistency
Optimization	Simulated annealing hyperparameters (temperature schedule, iteration count)
Sensitivity Scaling	Portfolio-size-dependent parameters adapting APY/TVL sensitivity

Remark 9. *Specific parameter values are calibrated through backtesting and production validation. The mathematical relationships governing parameter behavior are detailed in Sections 2–7. Contact the Sail team for deployment-specific configurations.*